House Price Prediction

July 22, 2023

[2]:
```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

[5]:
```
HouseDF = pd.read_csv('Housing_Price.csv')
HouseDF.head()
```

[5]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.45857	5.682861	7.009188	4.09	23086.80050	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701…
1	79248.64245	6.002900	6.730821	3.09	40173.07217	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA…
2	61287.06718	5.865890	8.512727	5.13	36882.15940	1.058988e+06	9127 Elizabeth Stravenue\nDanieltown, WI 06482…
3	63345.24005	7.188236	5.586729	3.26	34310.24283	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.19723	5.040555	7.839388	4.23	26354.10947	6.309435e+05	USNS Raymond\nFPO AE 09386

[6]:
```
HouseDF.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
 #   Column                        Non-Null Count  Dtype
---  ------                        --------------  -----
 0   Avg. Area Income              5000 non-null   float64
 1   Avg. Area House Age           5000 non-null   float64
```

```
 2   Avg. Area Number of Rooms     5000 non-null   float64
 3   Avg. Area Number of Bedrooms  5000 non-null   float64
 4   Area Population               5000 non-null   float64
 5   Price                         5000 non-null   float64
 6   Address                       5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

[7]:
```
HouseDF.describe()
```

[7]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
mean	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
std	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
min	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
25%	61480.562390	5.322283	6.299250	3.140000	29403.928700	9.975771e+05
50%	68804.286405	5.970429	7.002902	4.050000	36199.406690	1.232669e+06
75%	75783.338665	6.650808	7.665871	4.490000	42861.290770	1.471210e+06
max	107701.748400	9.519088	10.759588	6.500000	69621.713380	2.469066e+06

[8]:
```
HouseDF.columns
```

[8]:
```
Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
       'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
      dtype='object')
```

[9]:
```
sns.pairplot(HouseDF)
```

[9]:
```
<seaborn.axisgrid.PairGrid at 0x26c7239dff0>
```

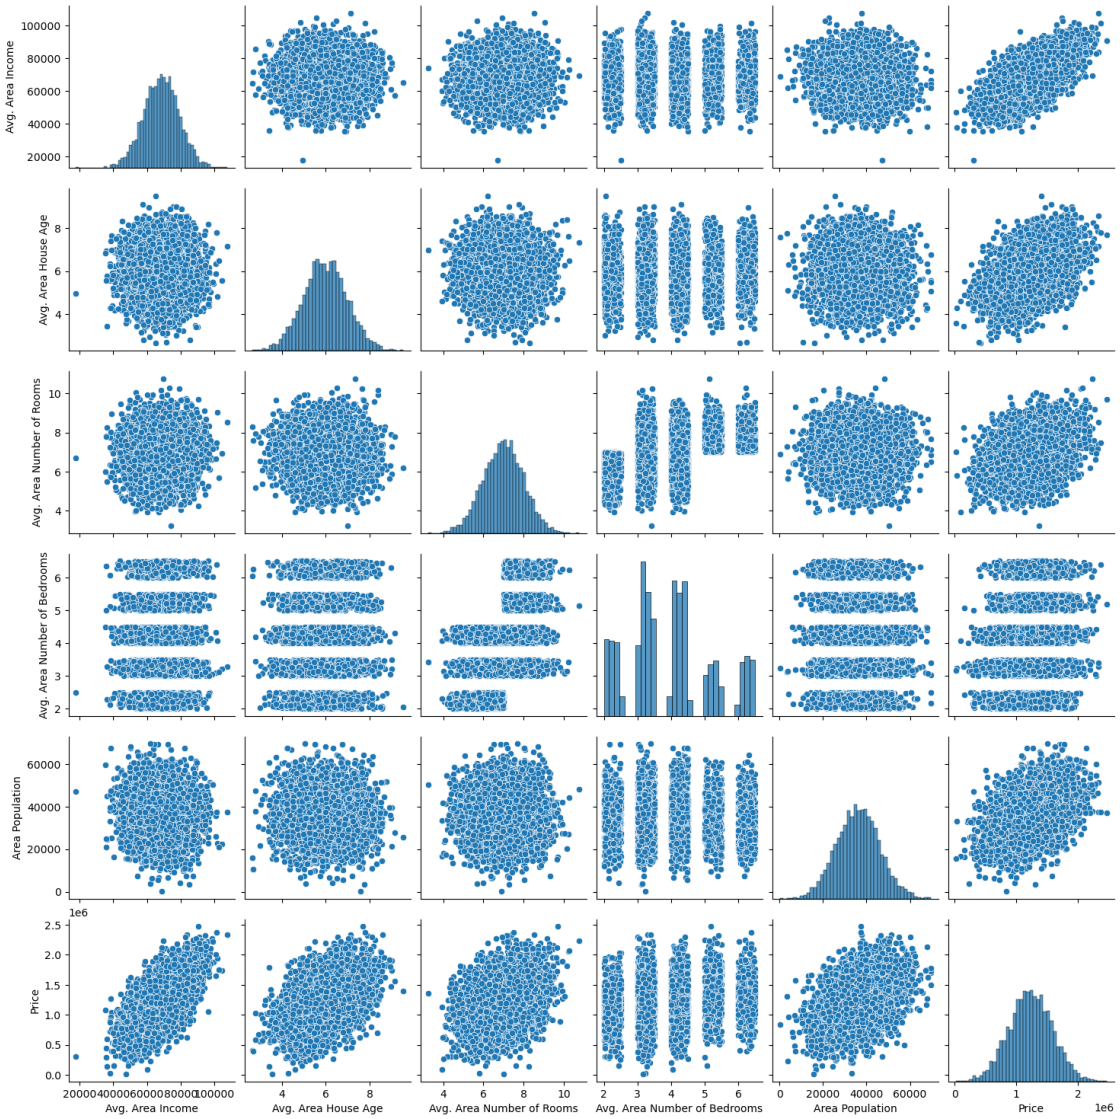

[36]:
```
sns.histplot(HouseDF['Price'])
```

[36]: <Axes: xlabel='Price', ylabel='Count'>

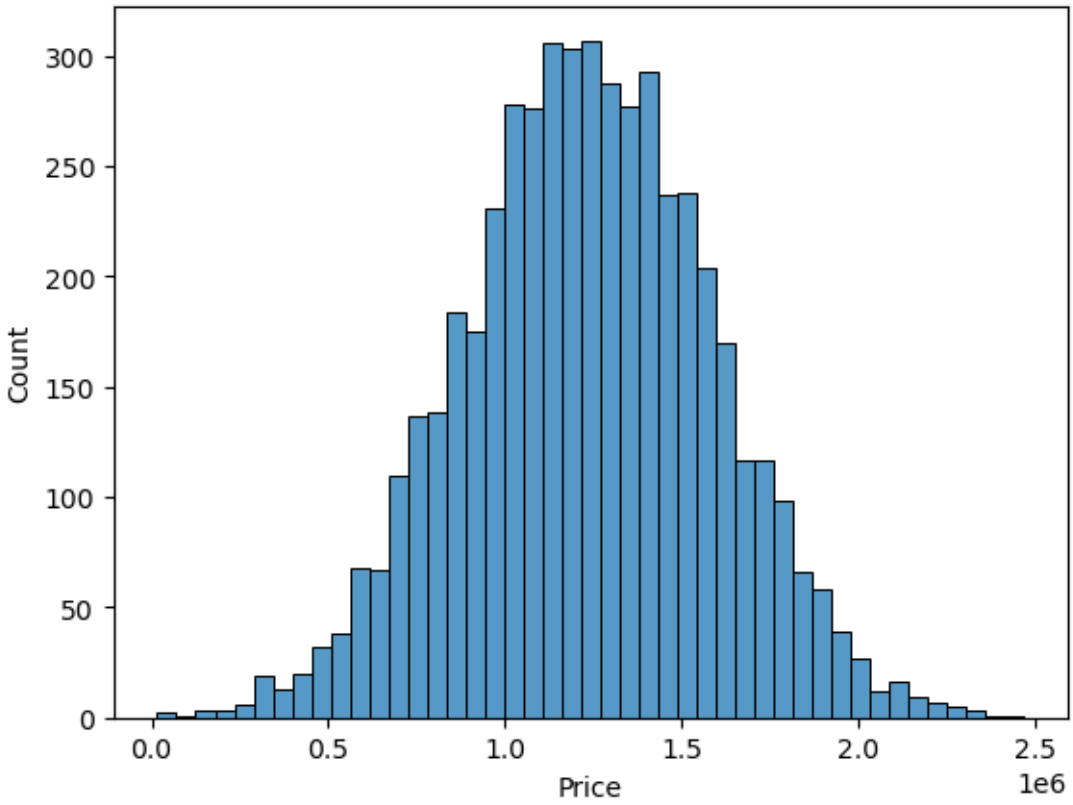

[38]:
```
X = HouseDF[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of␣
↪Rooms',
               'Avg. Area Number of Bedrooms', 'Area Population']]

y = HouseDF['Price']
```

[39]:
```
from sklearn.model_selection import train_test_split
```

[40]:
```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4,␣
↪random_state=101)
```

[41]:
```
from sklearn.linear_model import LinearRegression

lm = LinearRegression()

lm.fit(X_train,y_train)
```

[41]:
```
LinearRegression()
```

[42]:
```
print(lm.intercept_)
```

```
-2640159.7968132403
```

[52]:
```
coeff_DF = pd.DataFrame(lm.coef_,X.columns,columns=['Coefficient'])
coeff_DF
```

[52]:

	Coefficient
Avg. Area Income	21.528276
Avg. Area House Age	164883.282027
Avg. Area Number of Rooms	122368.678023
Avg. Area Number of Bedrooms	2233.801864
Area Population	15.150420

[53]:
```
predictions = lm.predict(X_test)
```

[54]:
```
plt.scatter(y_test,predictions)
```

[54]:
```
<matplotlib.collections.PathCollection at 0x26c7a2fbc10>
```

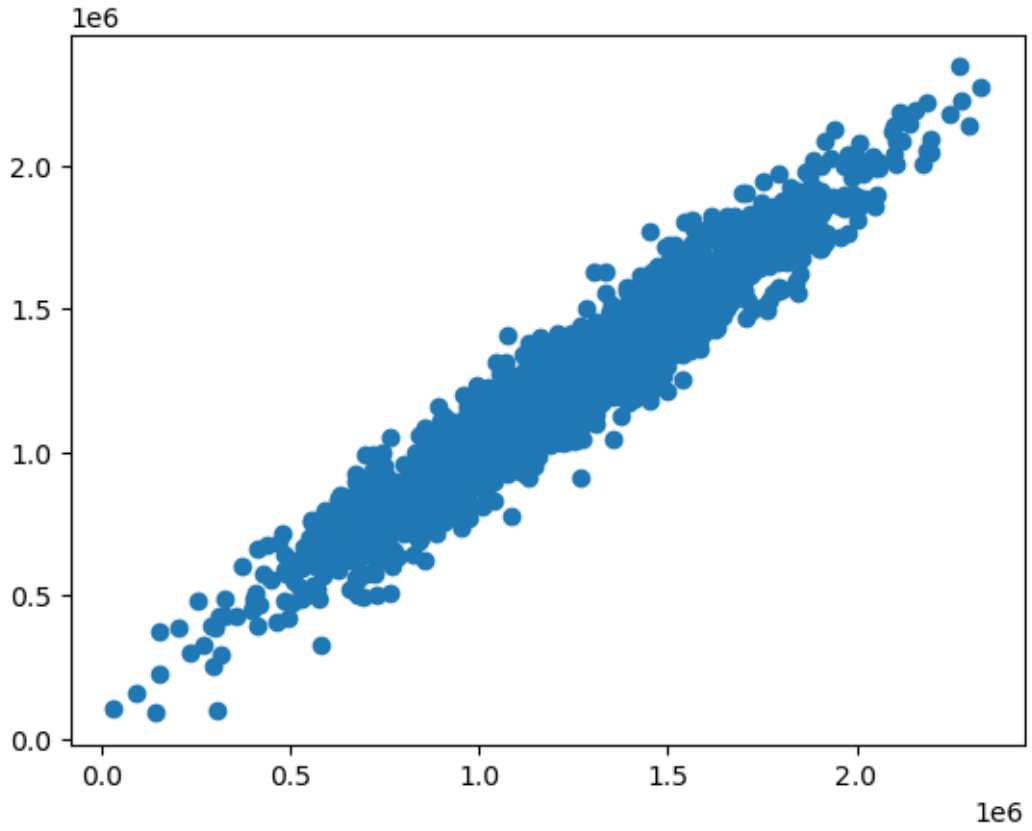

[59]:
```
sns.histplot((y_test,predictions),bins=50);
```

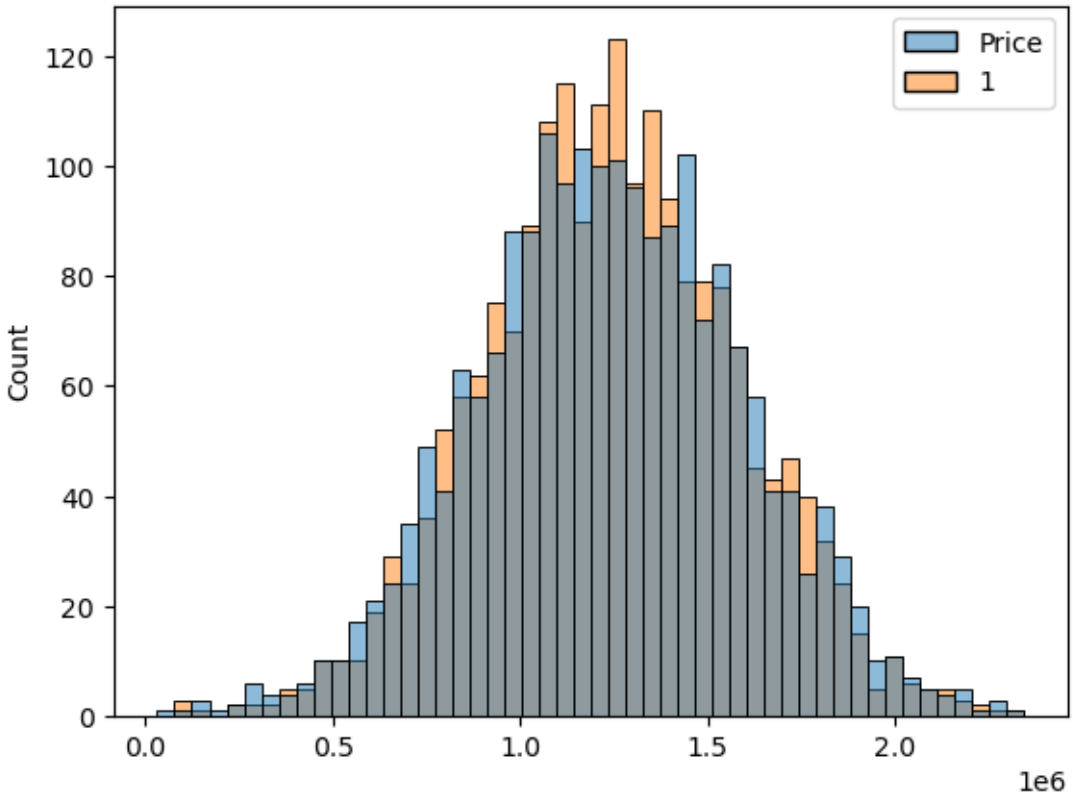

[60]:
```
from sklearn import metrics
```

[62]:
```
print('MAE:', metrics.mean_absolute_error(y_test, predictions))
print('MSE:', metrics.mean_squared_error(y_test, predictions))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

```
MAE: 82288.22250721784
MSE: 10460958905.775047
RMSE: 102278.82921589907
```

[ ]:
```

```